

Latency vs LLM Invocation Rate

Each point = one variant × scenario mean (31-run average)

